

0qgcyel1lo

April 30, 2024

```
[1]: import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
import os
import cv2
```

1 File paths for Images

```
[2]: train_data_path = 'C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//data//nиняcart_data//nиняcart_data//train'
test_data_path = 'C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//data//nиняcart_data//nиняcart_data//test'
```

```
[3]: os.listdir(train_data_path)
```

```
[3]: ['indian market', 'onion', 'potato', 'tomato']
```

2 Training image data count check

```
[4]: train_image_count_check = {}
for value in os.listdir(train_data_path):
    tmp_path = train_data_path+"//"+str(value)
    train_image_count_check[str(value)] = len(os.listdir(tmp_path))
train_image_count_check
```

```
[4]: {'indian market': 599, 'onion': 849, 'potato': 898, 'tomato': 789}
```

3 Testing image data count check

```
[5]: test_image_count_check = {}
for value in os.listdir(test_data_path):
    tmp_path = test_data_path+"//"+str(value)
```

```
test_image_count_check[str(value)] = len(os.listdir(tmp_path))
test_image_count_check
```

```
[5]: {'indian market': 81, 'onion': 83, 'potato': 81, 'tomato': 106}
```

4 Plotting Images from Training Folder

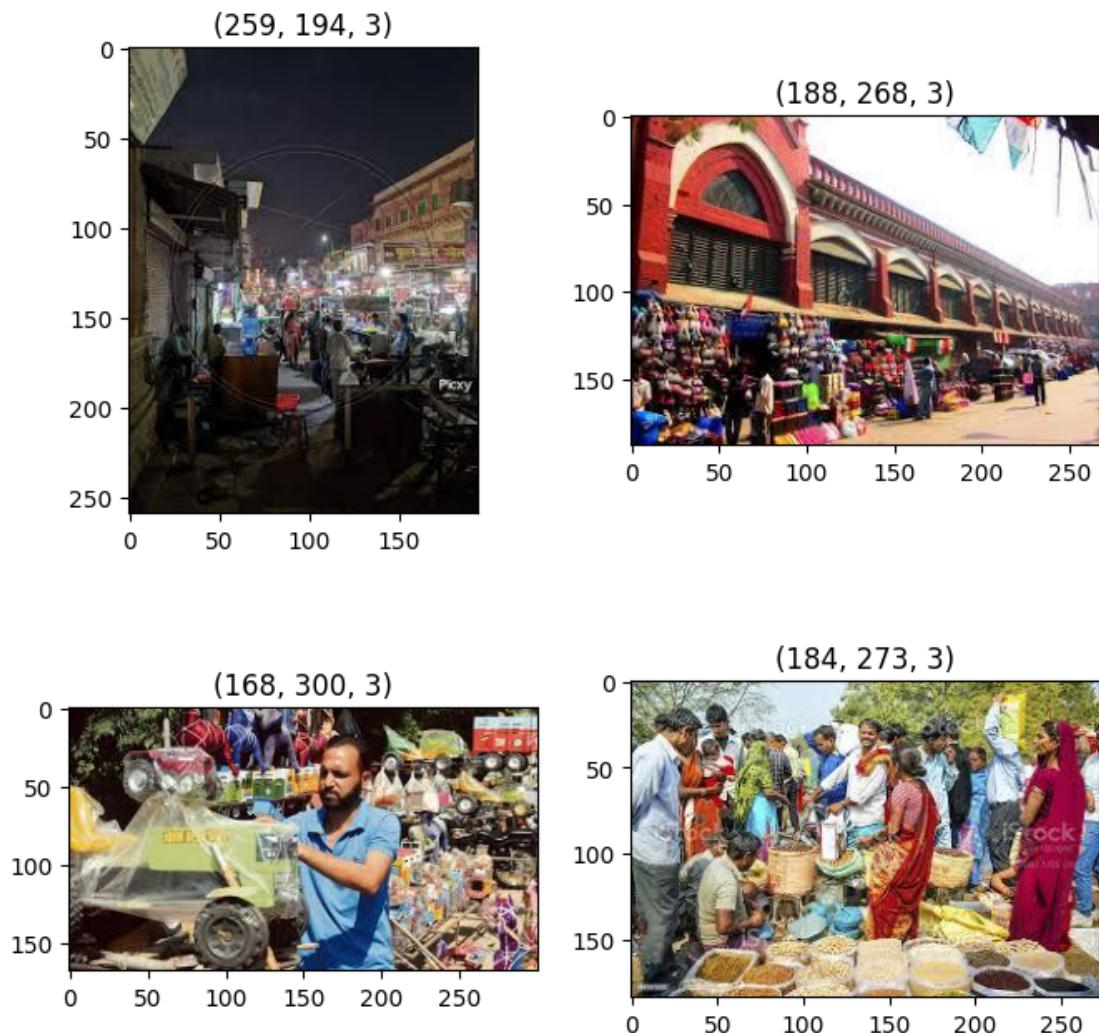
```
[6]: def plot_images(path,label,title):
    data = path+"//"+label
    first_four_images = os.listdir(data)[:4]
    fig, axs = plt.subplots(2,2, figsize=(8,8))
    for img,ax in enumerate(axs.flat):
        img_path = data+"//"+first_four_images[img]
        image = plt.imread(img_path)
        ax.set_title(image.shape)
        ax.imshow(image)
    fig.suptitle(title)
    plt.show()
```

```
[7]: os.listdir(train_data_path+"//"+'indian market')[-1:-4:-1]
```

```
[7]: ['village-market-827122__340.jpg', 'vendor-4049916__340.jpg', 'uyg.jpeg']
```

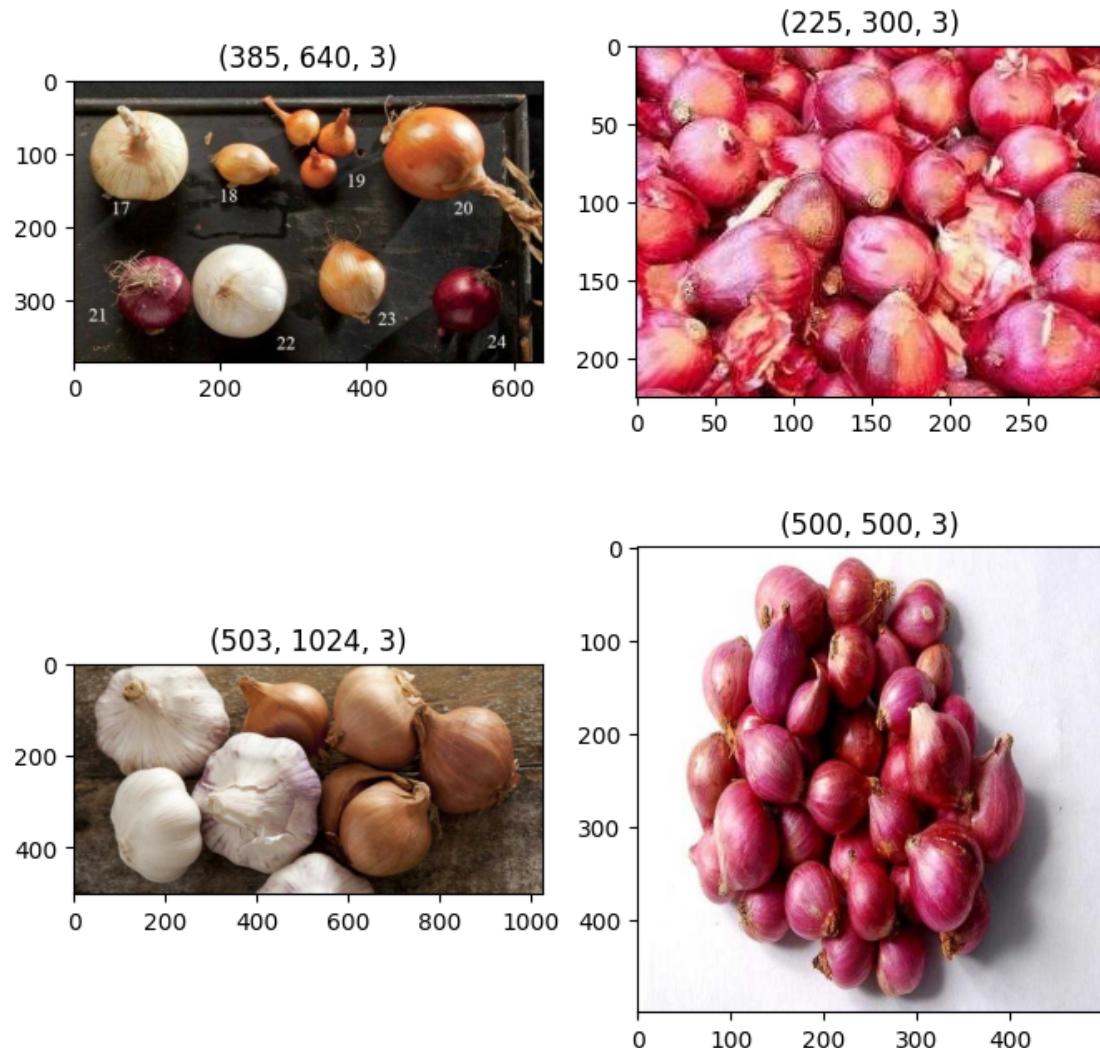
```
[8]: plot_images(train_data_path,"indian market","Images from the Indian Market")
```

Images from the Indian Market



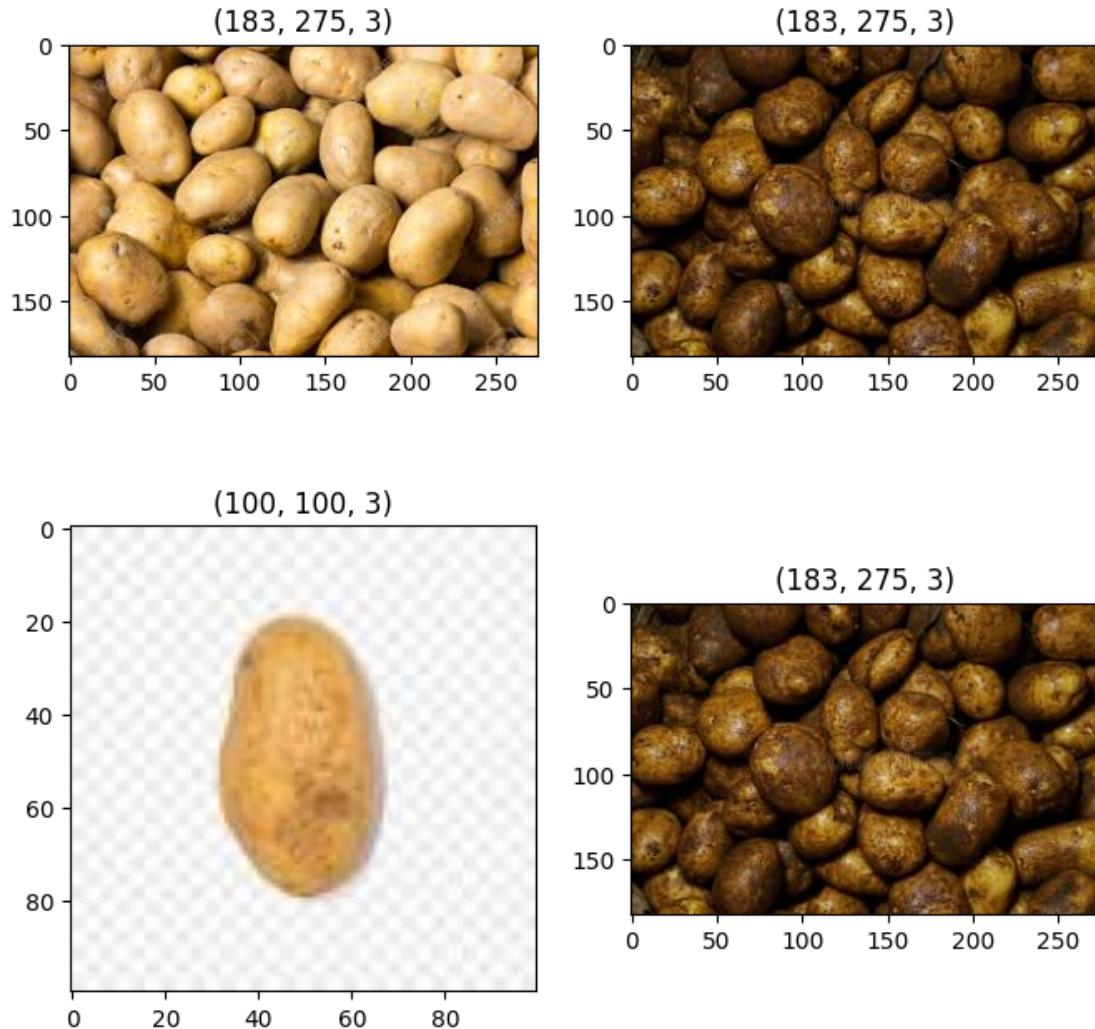
```
[9]: plot_images(train_data_path, "onion", "Images from the onion folder")
```

Images from the onion folder



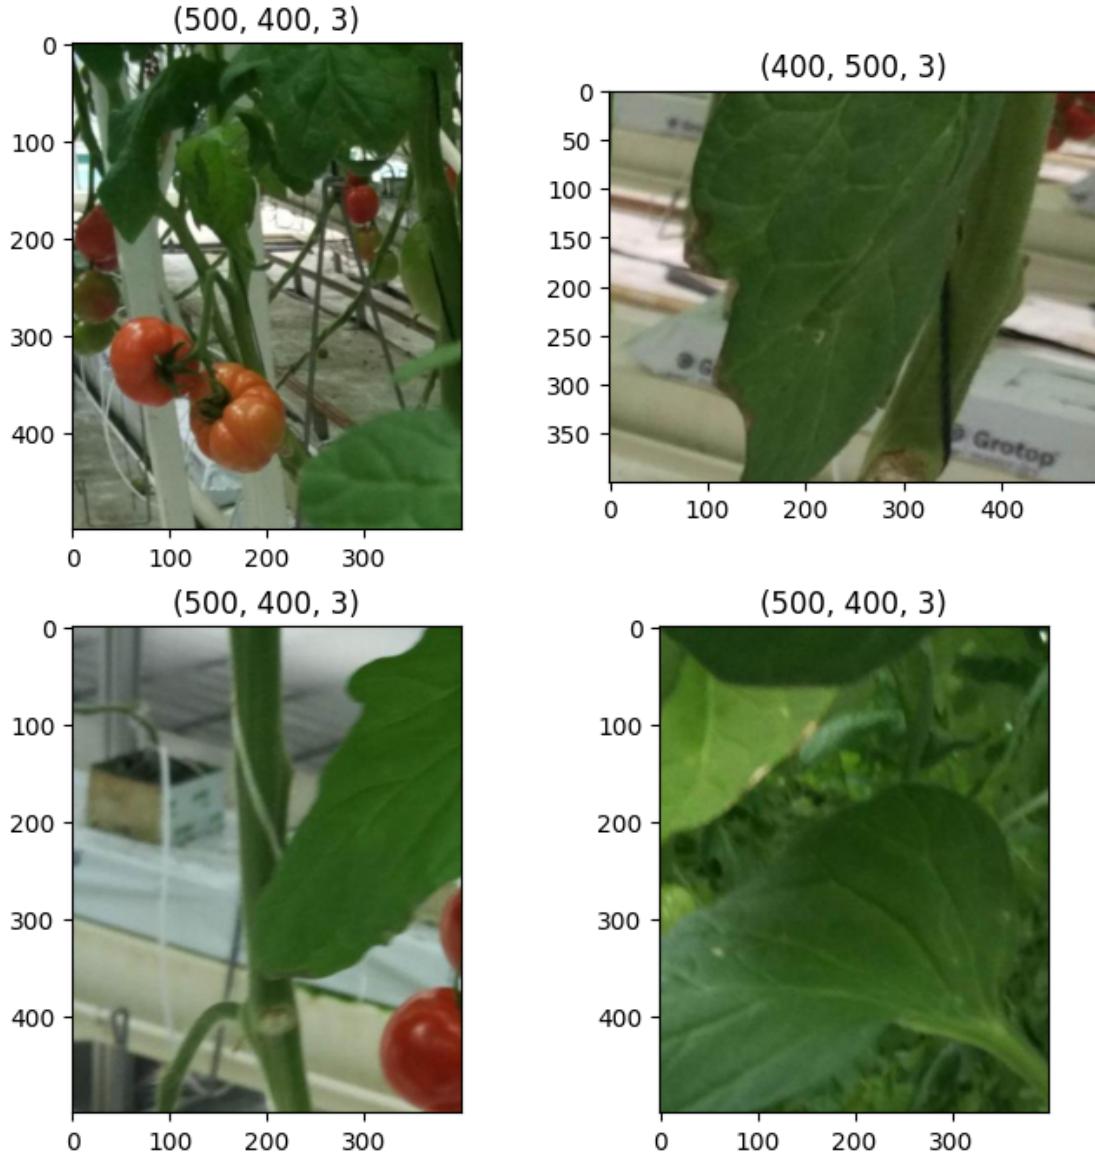
```
[10]: plot_images(train_data_path, "potato", "Images from the potato folder")
```

Images from the potato folder



```
[11]: plot_images(train_data_path, "tomato", "Images from the tomato folder")
```

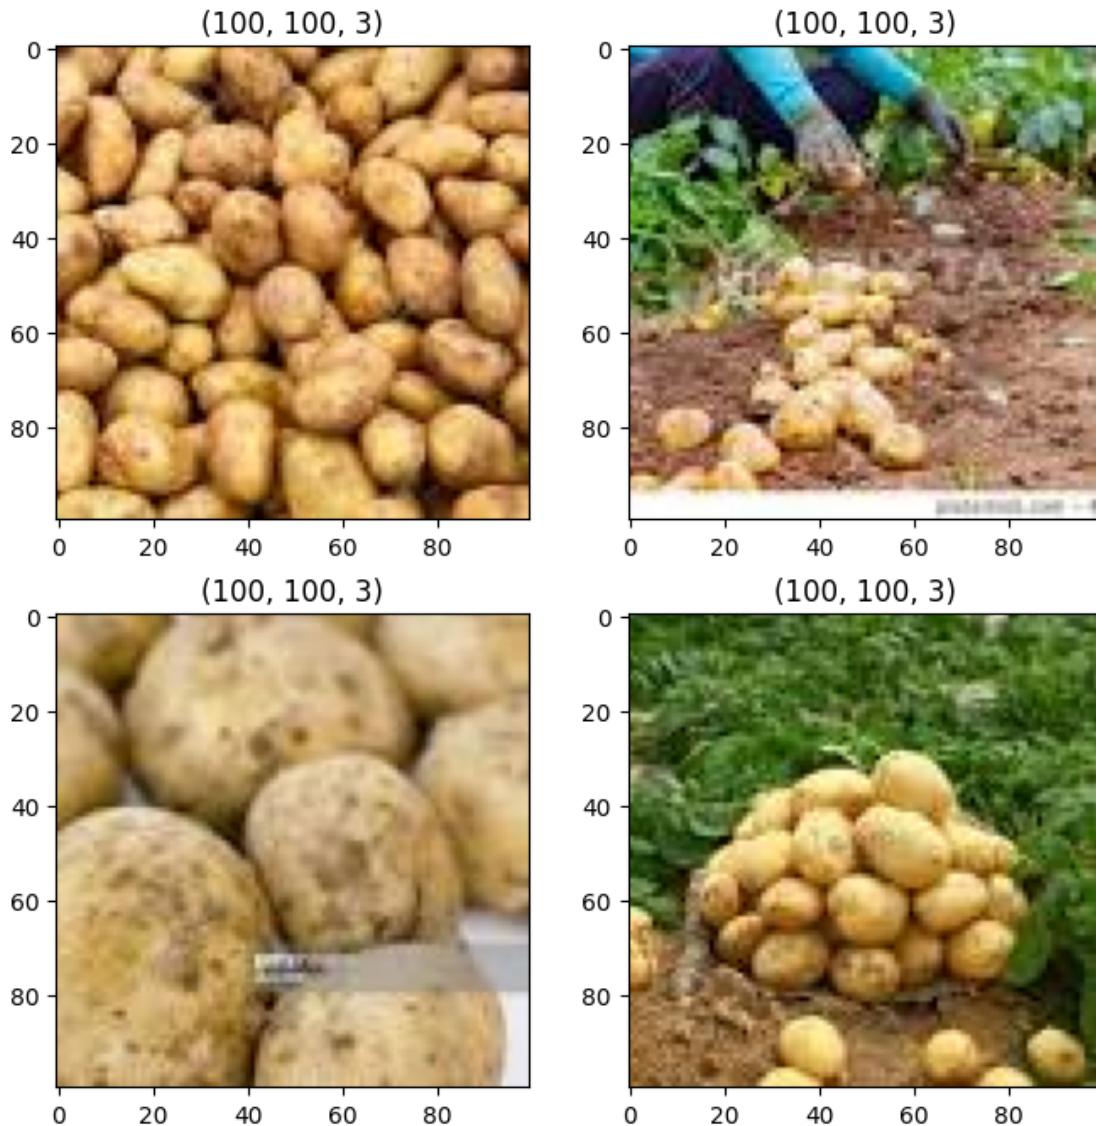
Images from the tomato folder



5 Ploting Images from Test Folder

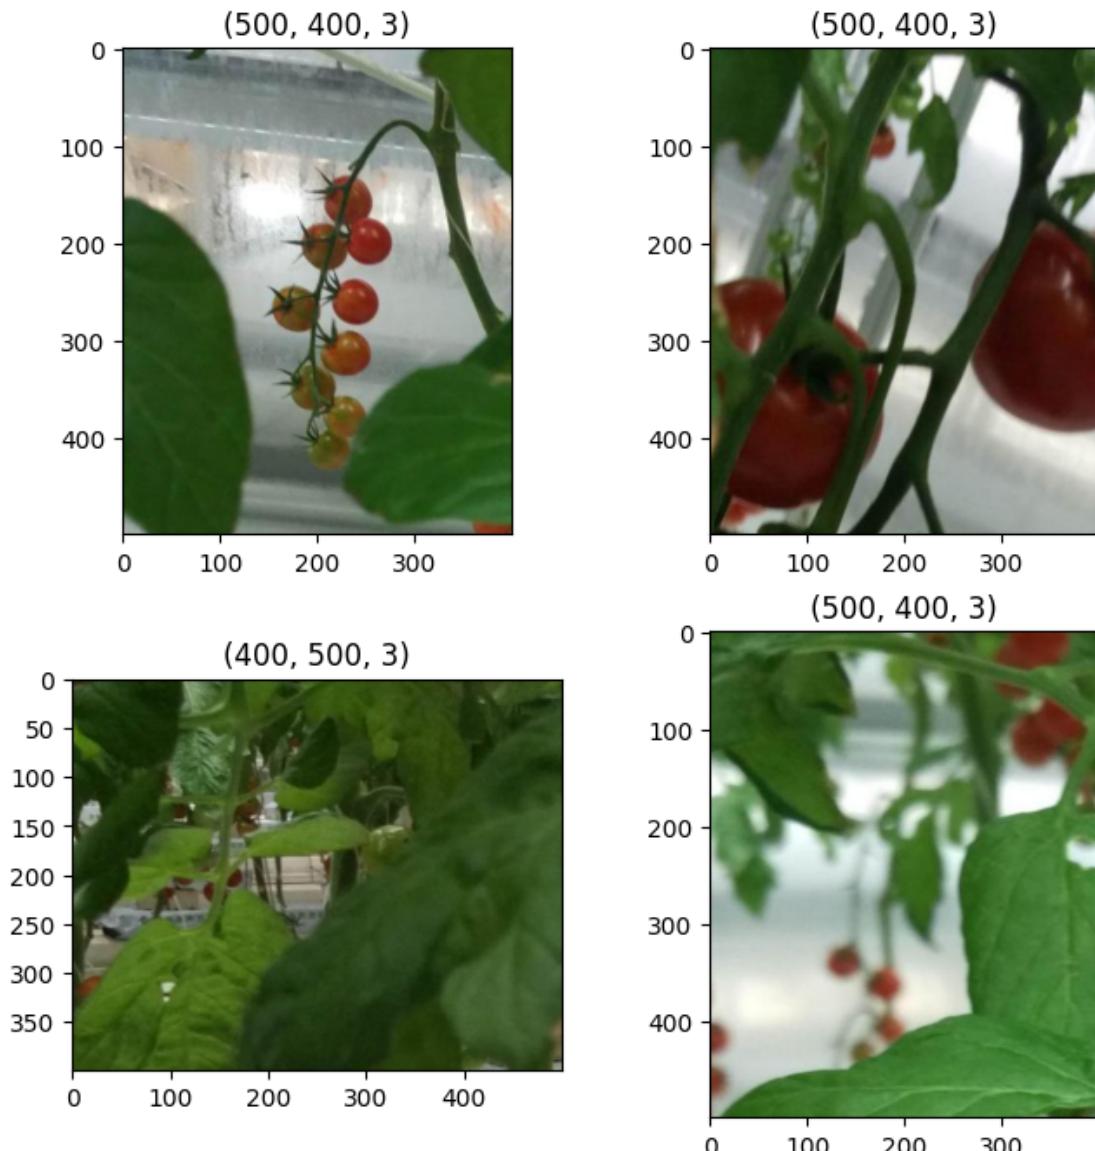
```
[12]: plot_images(test_data_path,"potato","Images from the potato Test folder")
```

Images from the potato Test folder



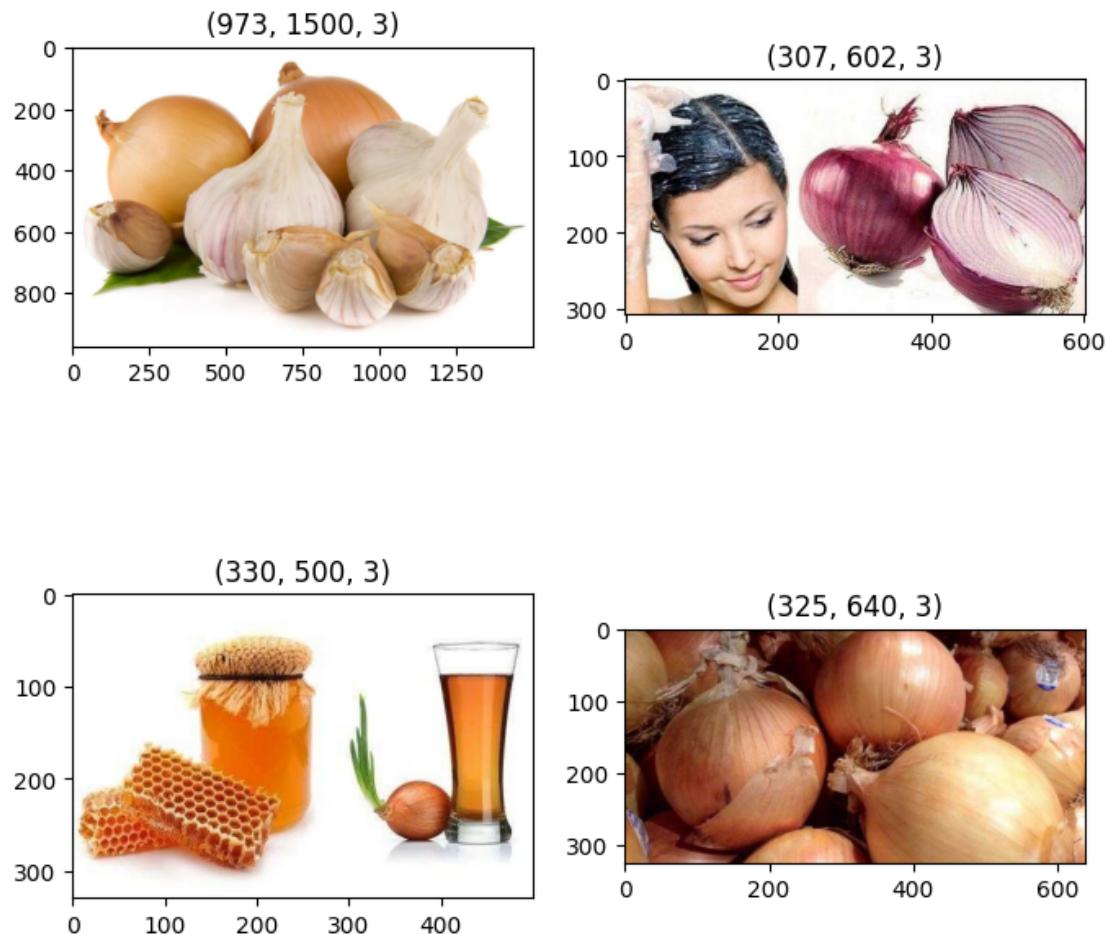
```
[13]: plot_images(test_data_path, "tomato", "Images from the tomato Test folder")
```

Images from the tomato Test folder



```
[14]: plot_images(test_data_path, "onion", "Images from the onion Test folder")
```

Images from the onion Test folder



Images from the indian market Test folder



6 Loading Training and Test Data

```
[16]: input_shape = (128,128)
train_data = tf.keras.utils.
    ↪image_dataset_from_directory(train_data_path,image_size=input_shape,batch_size=64,)
test_data = tf.keras.utils.
    ↪image_dataset_from_directory(test_data_path,image_size=input_shape,batch_size=64)
```

Found 3135 files belonging to 4 classes.

```
Found 351 files belonging to 4 classes.
```

```
[17]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import
    ↪Rescaling,RandomFlip,RandomTranslation,RandomRotation,RandomZoom,RandomContrast,RandomBright
```

7 Data Preprocessing and Data augmentation

```
[18]: preprocess_image = Sequential([
    Rescaling(scale=1./255)
])
```

```
[19]: train_data = train_data.map(lambda x,y : ↪
    ↪(preprocess_image(x,training=True),y),num_parallel_calls = tf.data.AUTOTUNE)
test_data = test_data.map(lambda x,y : ↪
    ↪(preprocess_image(x,training=True),y),num_parallel_calls = tf.data.AUTOTUNE)
```

```
[20]: len(train_data)
```

```
[20]: 49
```

```
[21]: # image_augmentation = Sequential([
#     RandomFlip("horizontal"),
#     RandomFlip("vertical"),
#     RandomRotation(factor=0.5),
#     RandomZoom(height_factor=-0.1,width_factor=-0.1)
# ])
```

```
[22]: # train_ds = train_data.map(lambda x,y : ↪
#     ↪(image_augmentation(x,training=True),y),num_parallel_calls =tf.data.
#     ↪AUTOTUNE)
```

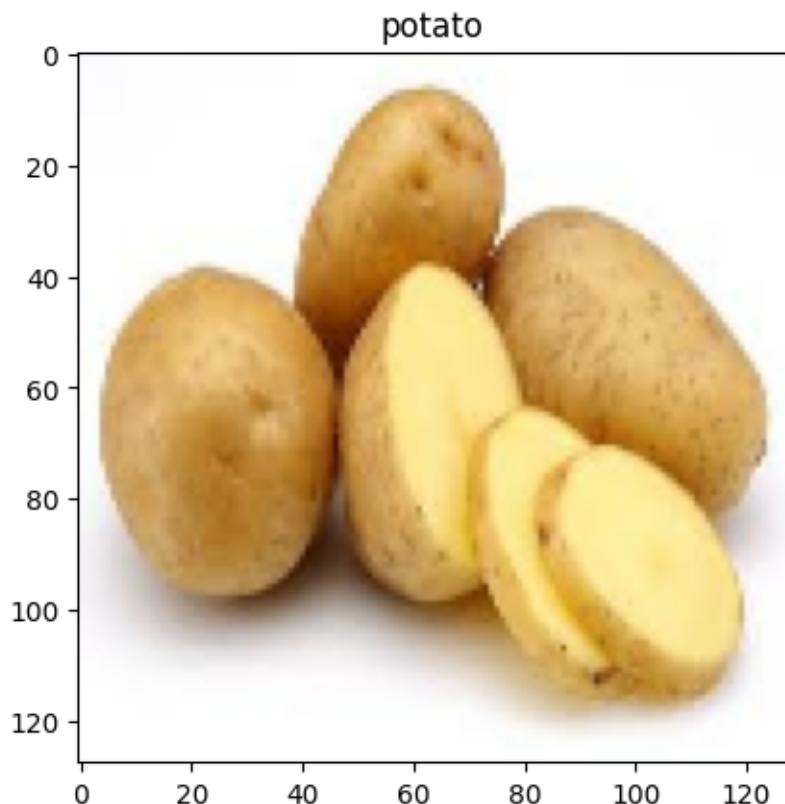
```
[23]: val_data = train_data.take(15)
train_ds = train_data.skip(15)
print(tf.data.experimental.cardinality(train_ds))
```

```
tf.Tensor(34, shape=(), dtype=int64)
```

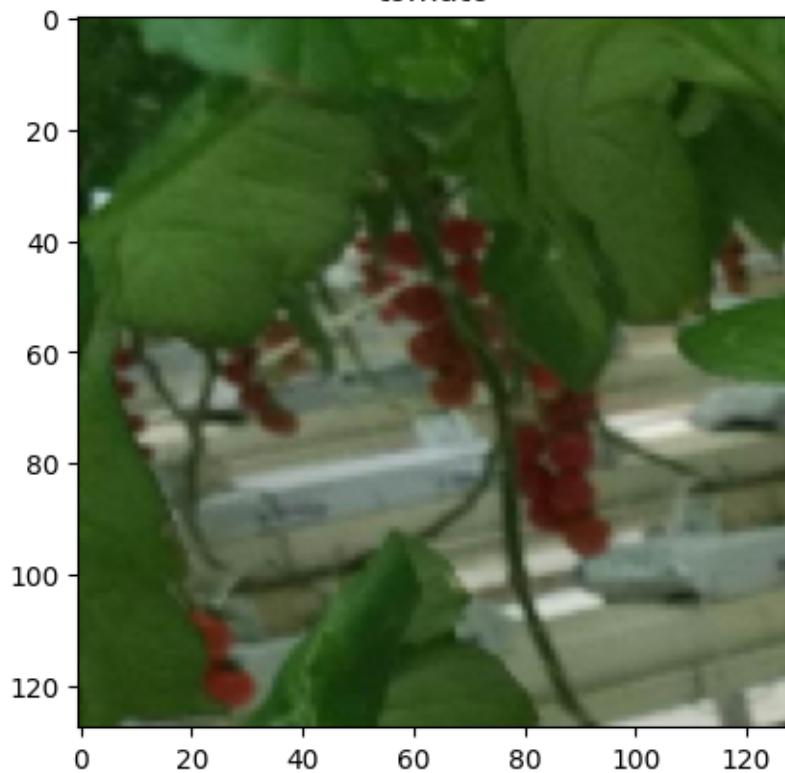
```
[24]: images,label = next(iter(train_ds))
```

```
[25]: reference_mapping_label = {
    2:'potato',
    3:'tomato',
    0:'indian market',
    1:'onion'
}
```

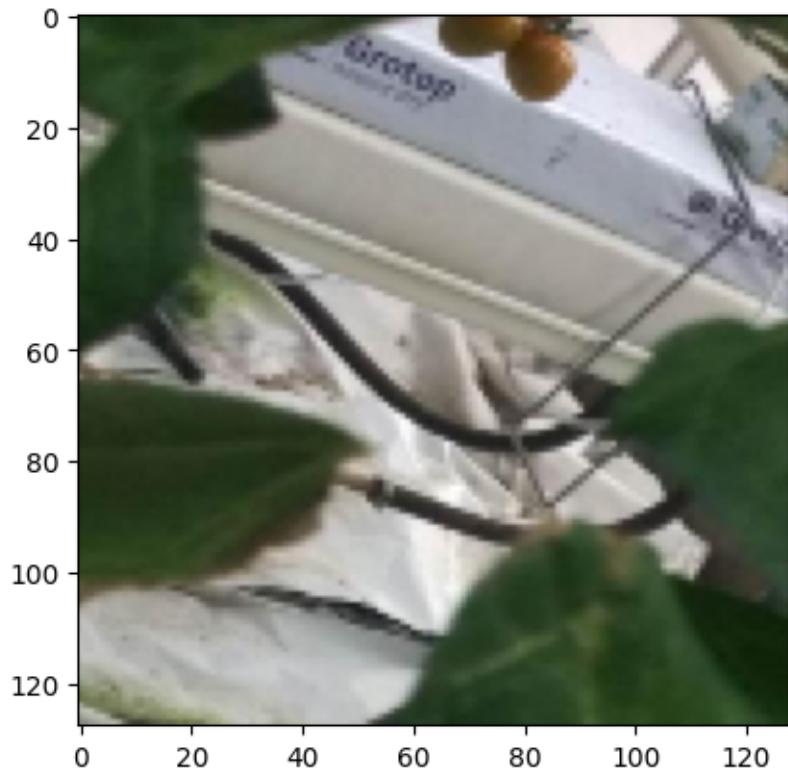
```
[26]: for img,label in zip(images.numpy()[:10],label.numpy()[:10]):  
    plt.title(reference_mapping_label[label])  
    # plt.title(label)  
    plt.imshow(img)  
    plt.show()
```



tomato

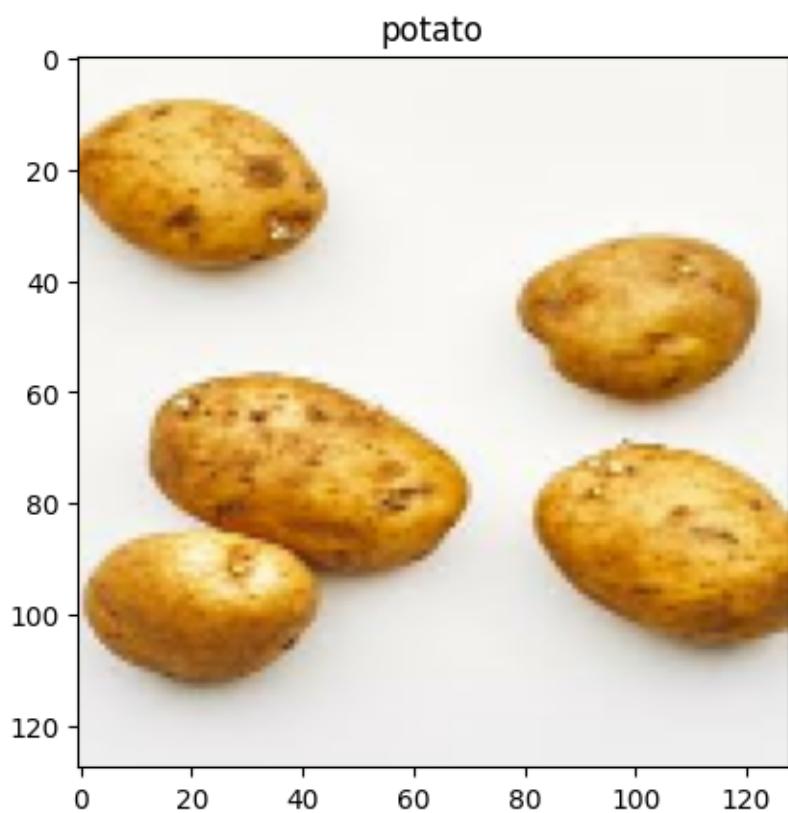


tomato



indian market





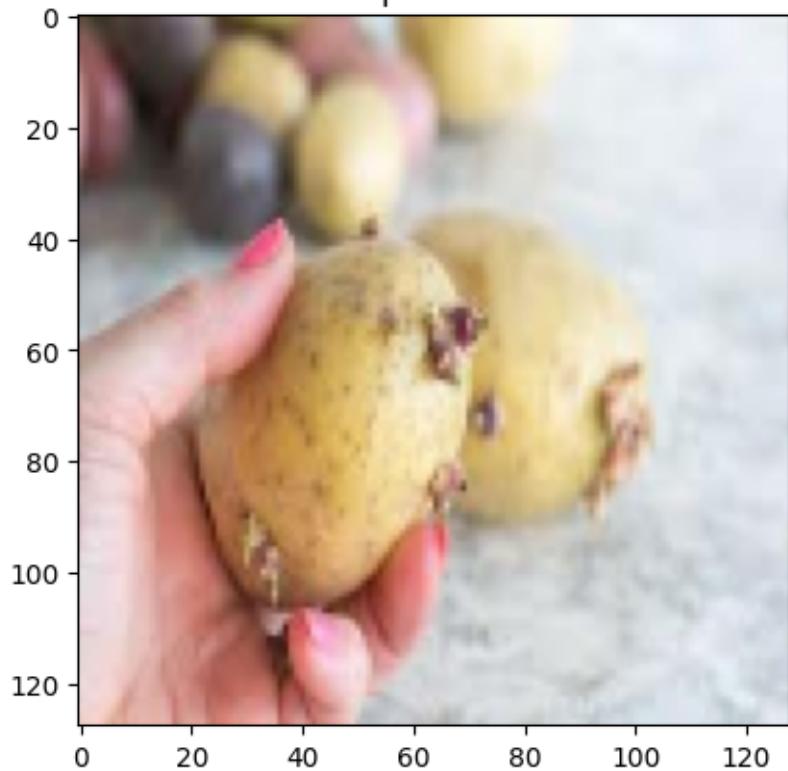
potato



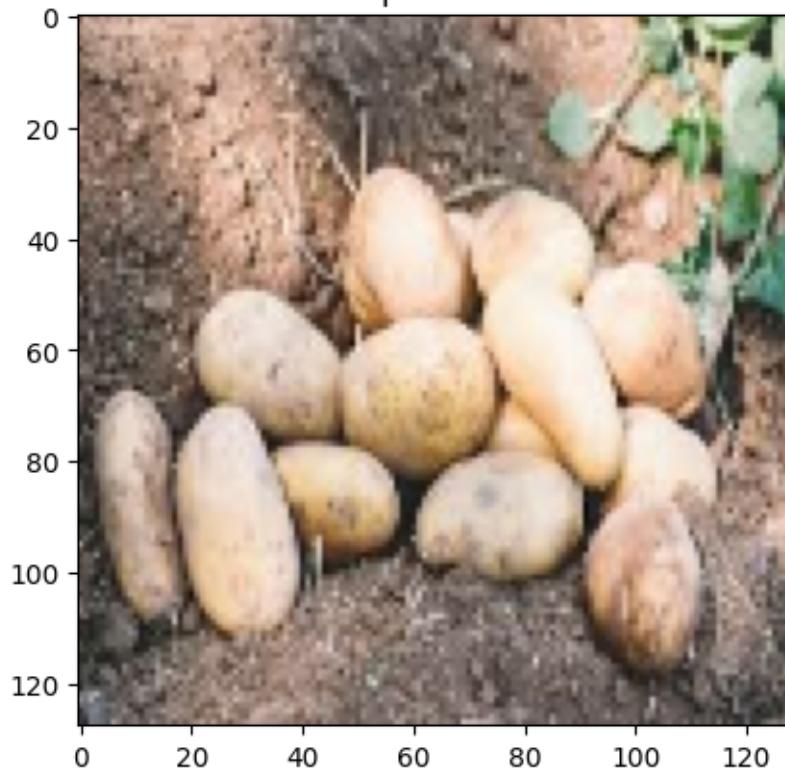
indian market

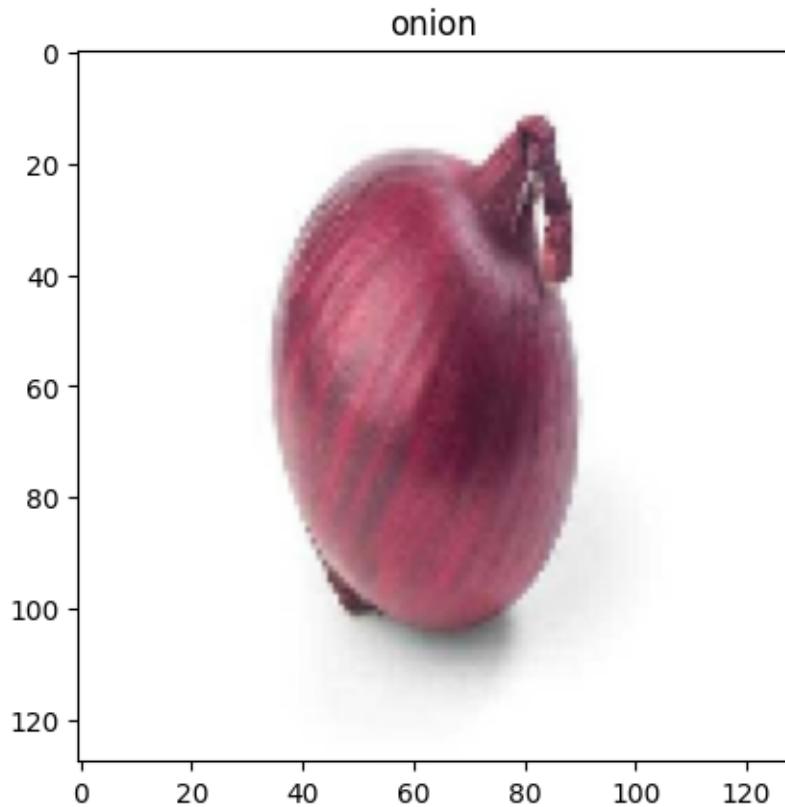


potato



potato





8 Scratch Implementation of CNN Model

```
[27]: from tensorflow.keras.layers import  
    Conv2D, Dropout, BatchNormalization, MaxPooling2D, GlobalAveragePooling2D, Dense  
from tensorflow.keras import regularizers
```

```
[28]: input_shape = (128, 128, 3)  
model = Sequential([  
  
    □  
    ↵Conv2D(filters=16, kernel_size=3, activation='relu', padding="same", input_shape=input_shape),  
        Conv2D(filters=16, kernel_size=3, activation='relu'),  
        Dropout(0.2),  
  
    □  
    ↵Conv2D(filters=32, kernel_size=3, activation='relu', padding="same", kernel_regularizer=regularizer  
    ↵L2(1e-4)),
```

```

    □
    ↵Conv2D(filters=32,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    MaxPooling2D(),
    Dropout(0.2),

    □
    ↵Conv2D(filters=64,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    □
    ↵Conv2D(filters=64,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    MaxPooling2D(),
    Dropout(0.3),

    □
    ↵Conv2D(filters=128,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    □
    ↵Conv2D(filters=128,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    MaxPooling2D(),
    Dropout(0.2),

    □
    ↵Conv2D(filters=256,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    □
    ↵Conv2D(filters=256,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    MaxPooling2D(),
    Dropout(0.2),

    □
    ↵Conv2D(filters=512,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    □
    ↵Conv2D(filters=512,kernel_size=3,activation='relu',padding="same",kernel_regularizer=regularizer,
    ↵L2(1e-4)),
    MaxPooling2D(),
    Dropout(0.2),

    GlobalAveragePooling2D(),

    Dense(units=4,activation='softmax')

```

])

[29]: model.summary()

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 16)	448
conv2d_1 (Conv2D)	(None, 126, 126, 16)	2320
dropout (Dropout)	(None, 126, 126, 16)	0
conv2d_2 (Conv2D)	(None, 126, 126, 32)	4640
conv2d_3 (Conv2D)	(None, 126, 126, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
dropout_1 (Dropout)	(None, 63, 63, 32)	0
conv2d_4 (Conv2D)	(None, 63, 63, 64)	18496
conv2d_5 (Conv2D)	(None, 63, 63, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 31, 31, 64)	0
dropout_2 (Dropout)	(None, 31, 31, 64)	0
conv2d_6 (Conv2D)	(None, 31, 31, 128)	73856
conv2d_7 (Conv2D)	(None, 31, 31, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 128)	0
dropout_3 (Dropout)	(None, 15, 15, 128)	0
conv2d_8 (Conv2D)	(None, 15, 15, 256)	295168
conv2d_9 (Conv2D)	(None, 15, 15, 256)	590080
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 256)	0

dropout_4 (Dropout)	(None, 7, 7, 256)	0
conv2d_10 (Conv2D)	(None, 7, 7, 512)	1180160
conv2d_11 (Conv2D)	(None, 7, 7, 512)	2359808
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 512)	0
dropout_5 (Dropout)	(None, 3, 3, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 4)	2052

=====

Total params: 4,720,788
Trainable params: 4,720,788
Non-trainable params: 0

```
[30]: from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
[31]: model.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])
filepath = "C://Users//gaura//Documents//GitHub//Scaler-Projects//ComputerVision//Project//Models//"
callbacks = [
    tf.keras.callbacks.ReduceLROnPlateau(
        monitor="val_loss", factor=0.3, patience=3, min_lr=0.00001
    ),
    tf.keras.callbacks.ModelCheckpoint(filepath, monitor='val_accuracy', mode='max', save_best_only=True),
    tf.keras.callbacks.EarlyStopping(
        monitor="val_loss", patience=3, min_delta=0.001, mode='min'
    )
]
```

```
[32]: tf.keras.backend.clear_session()
```

```
[33]: hist = model.fit(train_ds, validation_data=val_data, verbose=1, epochs=60, batch_size=32, callbacks=callbacks)
```

Epoch 1/60

```
34/34 [=====] - ETA: 0s - loss: 1.5214 - accuracy: 0.2602
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 48s 749ms/step - loss: 1.5214 - accuracy: 0.2602 - val_loss: 1.4708 - val_accuracy: 0.2677 - lr: 0.0010
Epoch 2/60
34/34 [=====] - 13s 322ms/step - loss: 1.4558 - accuracy: 0.2717 - val_loss: 1.4456 - val_accuracy: 0.2646 - lr: 0.0010
Epoch 3/60
34/34 [=====] - ETA: 0s - loss: 1.4217 - accuracy: 0.3149
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 16s 419ms/step - loss: 1.4217 - accuracy: 0.3149 - val_loss: 1.3761 - val_accuracy: 0.4187 - lr: 0.0010
Epoch 4/60
34/34 [=====] - ETA: 0s - loss: 1.1786 - accuracy: 0.4782
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 16s 427ms/step - loss: 1.1786 - accuracy: 0.4782 - val_loss: 0.9694 - val_accuracy: 0.5344 - lr: 0.0010
Epoch 5/60
```

```
34/34 [=====] - ETA: 0s - loss: 0.9562 - accuracy: 0.5375
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 16s 434ms/step - loss: 0.9562 - accuracy: 0.5375 - val_loss: 1.0074 - val_accuracy: 0.5698 - lr: 0.0010
Epoch 6/60
34/34 [=====] - ETA: 0s - loss: 0.8803 - accuracy: 0.6106
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 16s 423ms/step - loss: 0.8803 - accuracy: 0.6106 - val_loss: 0.8137 - val_accuracy: 0.6396 - lr: 0.0010
Epoch 7/60
34/34 [=====] - 12s 323ms/step - loss: 0.8731 - accuracy: 0.6078 - val_loss: 0.8669 - val_accuracy: 0.5948 - lr: 0.0010
Epoch 8/60
34/34 [=====] - ETA: 0s - loss: 0.7927 - accuracy: 0.6317
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 15s 424ms/step - loss: 0.7927 - accuracy: 0.6317 - val_loss: 0.7682 - val_accuracy: 0.6438 - lr: 0.0010
Epoch 9/60
```

```
34/34 [=====] - ETA: 0s - loss: 0.8032 - accuracy: 0.6487
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 15s 413ms/step - loss: 0.8032 - accuracy: 0.6487 - val_loss: 0.7193 - val_accuracy: 0.6865 - lr: 0.0010
Epoch 10/60
34/34 [=====] - 12s 318ms/step - loss: 0.7319 - accuracy: 0.6874 - val_loss: 0.7117 - val_accuracy: 0.6750 - lr: 0.0010
Epoch 11/60
34/34 [=====] - 12s 319ms/step - loss: 0.7138 - accuracy: 0.6947 - val_loss: 0.7320 - val_accuracy: 0.6635 - lr: 0.0010
Epoch 12/60
34/34 [=====] - ETA: 0s - loss: 0.6973 - accuracy: 0.7140
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
34/34 [=====] - 15s 422ms/step - loss: 0.6973 - accuracy: 0.7140 - val_loss: 0.6503 - val_accuracy: 0.7240 - lr: 0.0010
Epoch 13/60
34/34 [=====] - 12s 321ms/step - loss: 0.6503 - accuracy: 0.7361 - val_loss: 0.6361 - val_accuracy: 0.7146 - lr: 0.0010
Epoch 14/60
34/34 [=====] - ETA: 0s - loss: 0.7298 - accuracy: 0.6791
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 12). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer Vision//Project//Models\\assets
```

```
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-  
Projects//Computer Vision//Project//Models\\assets  
  
34/34 [=====] - 16s 445ms/step - loss: 0.7298 -  
accuracy: 0.6791 - val_loss: 0.6884 - val_accuracy: 0.7302 - lr: 0.0010  
Epoch 15/60  
34/34 [=====] - 12s 325ms/step - loss: 0.7228 -  
accuracy: 0.6952 - val_loss: 0.7227 - val_accuracy: 0.6573 - lr: 0.0010  
Epoch 16/60  
34/34 [=====] - ETA: 0s - loss: 0.6418 - accuracy:  
0.7434  
  
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,  
_jit_compiled_convolution_op, _jit_compiled_convolution_op,  
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing  
5 of 12). These functions will not be directly callable after loading.  
  
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-  
Projects//Computer Vision//Project//Models\\assets  
  
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-  
Projects//Computer Vision//Project//Models\\assets  
  
34/34 [=====] - 16s 424ms/step - loss: 0.6418 -  
accuracy: 0.7434 - val_loss: 0.5810 - val_accuracy: 0.7688 - lr: 0.0010  
Epoch 17/60  
34/34 [=====] - 12s 322ms/step - loss: 0.5671 -  
accuracy: 0.7770 - val_loss: 0.5846 - val_accuracy: 0.7688 - lr: 0.0010  
Epoch 18/60  
34/34 [=====] - 12s 321ms/step - loss: 0.6186 -  
accuracy: 0.7655 - val_loss: 0.6650 - val_accuracy: 0.7115 - lr: 0.0010  
Epoch 19/60  
34/34 [=====] - 12s 319ms/step - loss: 0.5816 -  
accuracy: 0.7623 - val_loss: 0.5665 - val_accuracy: 0.7656 - lr: 0.0010  
Epoch 20/60  
34/34 [=====] - ETA: 0s - loss: 0.5385 - accuracy:  
0.7926  
  
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,  
_jit_compiled_convolution_op, _jit_compiled_convolution_op,  
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing  
5 of 12). These functions will not be directly callable after loading.  
  
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-  
Projects//Computer Vision//Project//Models\\assets  
  
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-  
Projects//Computer Vision//Project//Models\\assets  
  
34/34 [=====] - 16s 428ms/step - loss: 0.5385 -  
accuracy: 0.7926 - val_loss: 0.5301 - val_accuracy: 0.7865 - lr: 0.0010  
Epoch 21/60
```

```

34/34 [=====] - 13s 325ms/step - loss: 0.5324 -
accuracy: 0.7949 - val_loss: 0.5746 - val_accuracy: 0.7635 - lr: 0.0010
Epoch 22/60
34/34 [=====] - ETA: 0s - loss: 0.5277 - accuracy:
0.7954

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing
5 of 12). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-
Projects//Computer Vision//Project//Models\\assets

INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-
Projects//Computer Vision//Project//Models\\assets

34/34 [=====] - 16s 436ms/step - loss: 0.5277 -
accuracy: 0.7954 - val_loss: 0.5222 - val_accuracy: 0.7969 - lr: 0.0010
Epoch 23/60
34/34 [=====] - 12s 324ms/step - loss: 0.5017 -
accuracy: 0.8138 - val_loss: 0.5512 - val_accuracy: 0.7760 - lr: 0.0010
Epoch 24/60
34/34 [=====] - ETA: 0s - loss: 0.4697 - accuracy:
0.8166

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing
5 of 12). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-
Projects//Computer Vision//Project//Models\\assets

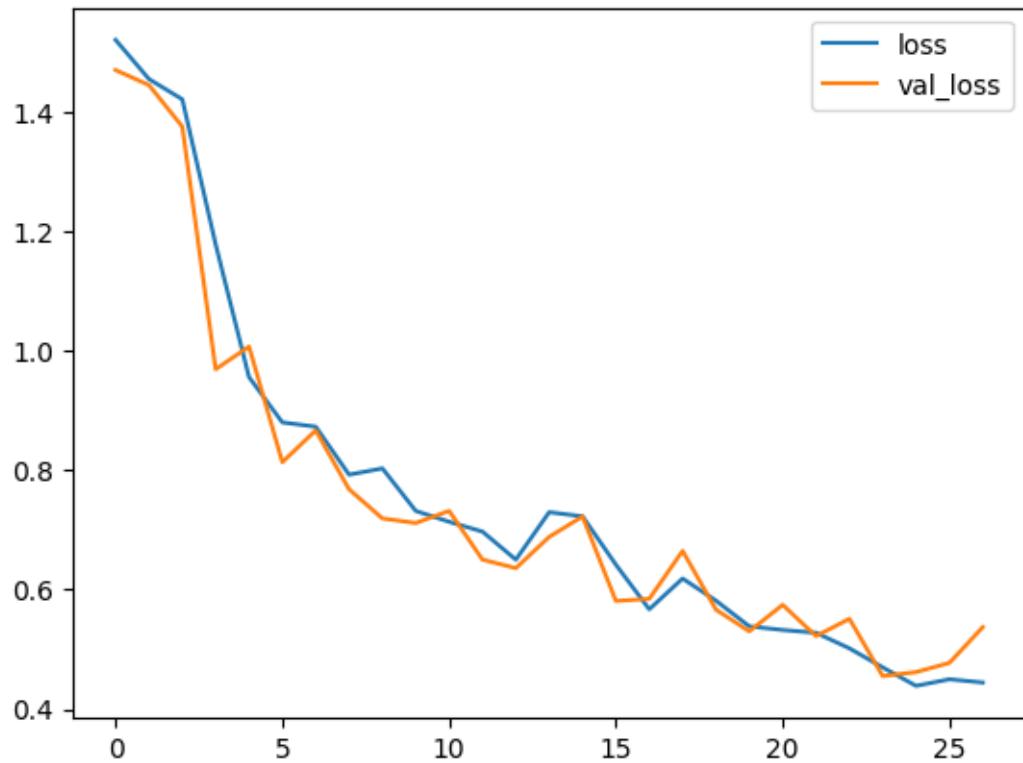
INFO:tensorflow:Assets written to: C://Users//gaura//Documents//GitHub//Scaler-
Projects//Computer Vision//Project//Models\\assets

34/34 [=====] - 15s 425ms/step - loss: 0.4697 -
accuracy: 0.8166 - val_loss: 0.4552 - val_accuracy: 0.8229 - lr: 0.0010
Epoch 25/60
34/34 [=====] - 12s 324ms/step - loss: 0.4388 -
accuracy: 0.8271 - val_loss: 0.4619 - val_accuracy: 0.8188 - lr: 0.0010
Epoch 26/60
34/34 [=====] - 12s 323ms/step - loss: 0.4498 -
accuracy: 0.8230 - val_loss: 0.4773 - val_accuracy: 0.8052 - lr: 0.0010
Epoch 27/60
34/34 [=====] - 12s 323ms/step - loss: 0.4444 -
accuracy: 0.8221 - val_loss: 0.5372 - val_accuracy: 0.8156 - lr: 0.0010

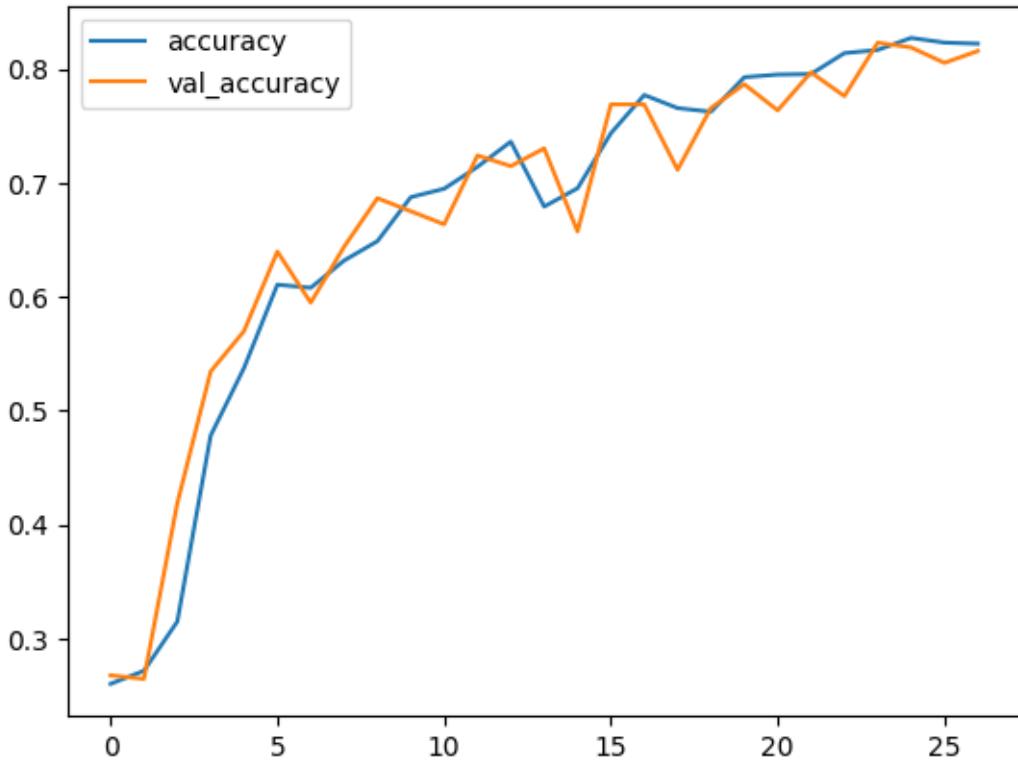
```

[34]: plt.plot(hist.history['loss'],label='loss')
plt.plot(hist.history['val_loss'],label='val_loss')

```
plt.legend()  
plt.show()
```



```
[35]: plt.plot(hist.history['accuracy'],label='accuracy')  
plt.plot(hist.history['val_accuracy'],label='val_accuracy')  
plt.legend()  
plt.show()
```



```
[36]: load_model = tf.keras.models.load_model('C://Users//gaura//Documents//GitHub//  
Scaler-Projects//Computer Vision//Project//Models')
```

```
[37]: # Evaluate the model  
loss, acc = load_model.evaluate(test_data, verbose=2)  
print("Restored model, accuracy: {:.2f}%".format(100 * acc))
```

6/6 - 4s - loss: 0.5625 - accuracy: 0.7521 - 4s/epoch - 600ms/step
Restored model, accuracy: 75.21%

```
[38]: y_pred = load_model.predict(test_data)
```

6/6 [=====] - 1s 55ms/step

```
[39]: pred = np.argmax(y_pred, axis=1)
```

```
[40]: ytest = []  
for img,label in test_data:  
    labels = label.numpy()  
    ytest.append(labels)
```

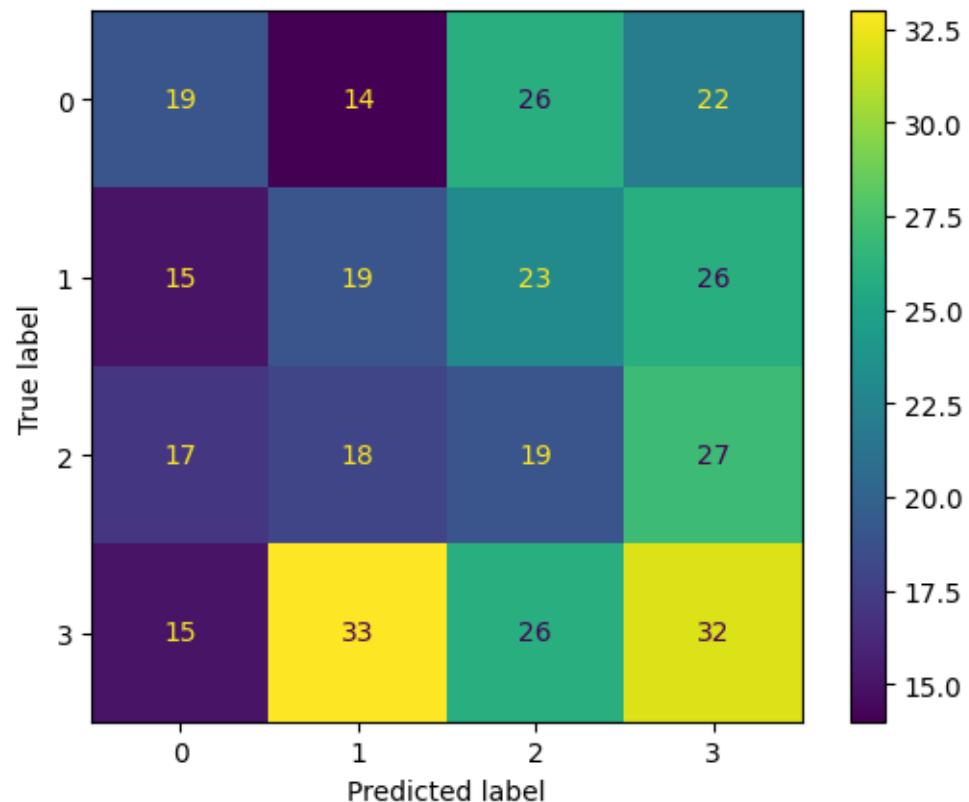
```
[41]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
```

```
[42]: print(classification_report(ytest,pred))
```

	precision	recall	f1-score	support
0	0.29	0.23	0.26	81
1	0.23	0.23	0.23	83
2	0.20	0.23	0.22	81
3	0.30	0.30	0.30	106
accuracy			0.25	351
macro avg	0.25	0.25	0.25	351
weighted avg	0.26	0.25	0.25	351

```
[43]: ConfusionMatrixDisplay(confusion_matrix(ytest,pred)).plot()
```

```
[43]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x225772bfef0>
```



```
[44]: images,label = next(iter(test_data))
images = images.numpy()
label = label.numpy()
```

```
[45]: reference_mapping_label = {
    2:'potato',
    3:'tomato',
    0:'indian market',
    1:'onion'
}
```

```
[46]: CLASS_NAMES = ['indian market','onion','potato','tomato']
```

```
[47]: for image, label in test_data.take(2):
    plt.imshow(image[2])
    plt.title(CLASS_NAMES[label[2].numpy()])
    plt.axis('off')
prediction = load_model.predict(test_data.take(2))[2]
prediction = np.argmax(prediction)
#print(prediction)
scores = [1 - prediction, prediction]
for score, name in zip(scores, CLASS_NAMES):
    print("This image is %.2f percent %s" % ((100 * score), name))
%time
```

2/2 [=====] - 0s 149ms/step

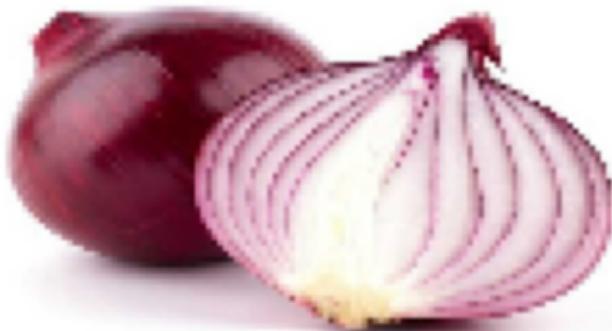
This image is -200.00 percent indian market

This image is 300.00 percent onion

CPU times: total: 0 ns

Wall time: 0 ns

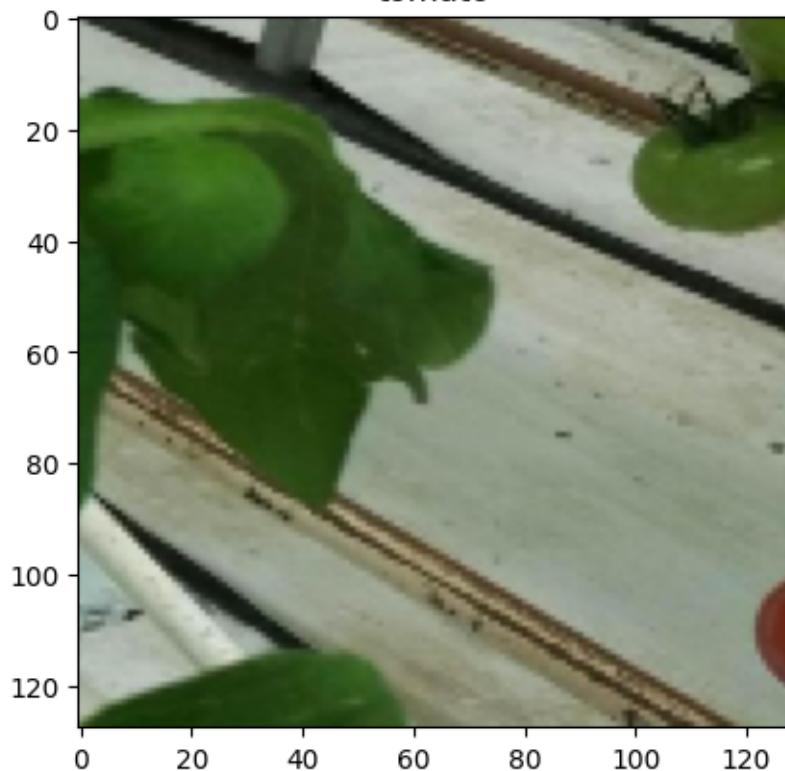
onion



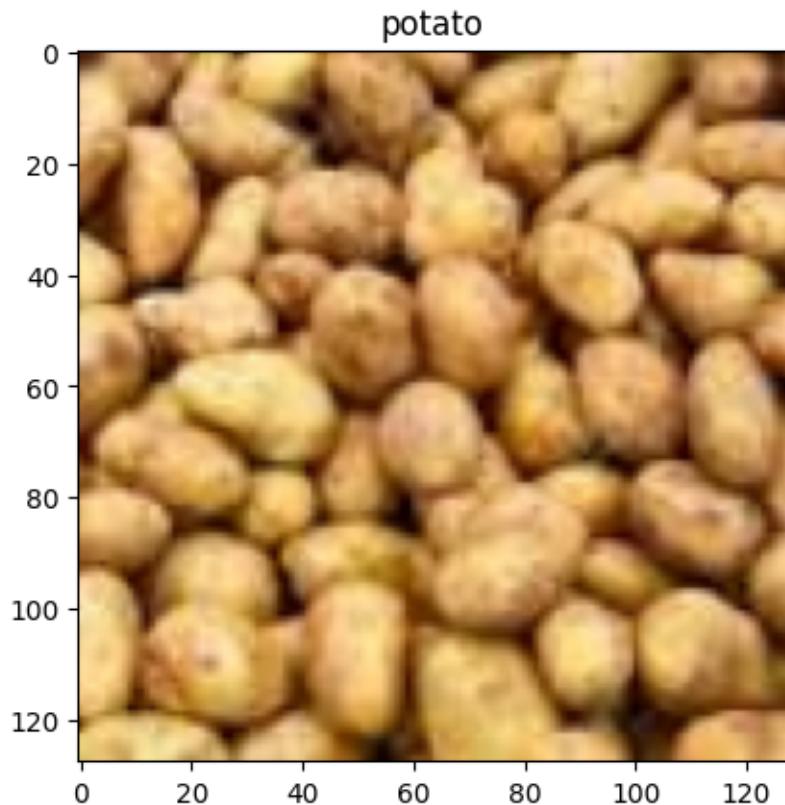
```
[48]: for img in images[:5]:
    #print(img.shape)
    temp = (np.expand_dims(img, axis=0))
    pred = load_model.predict(temp)
    labels = reference_mapping_label[np.argmax(pred)]
    plt.imshow(img)
    plt.title(labels)
    plt.show()
```

1/1 [=====] - 1s 1s/step

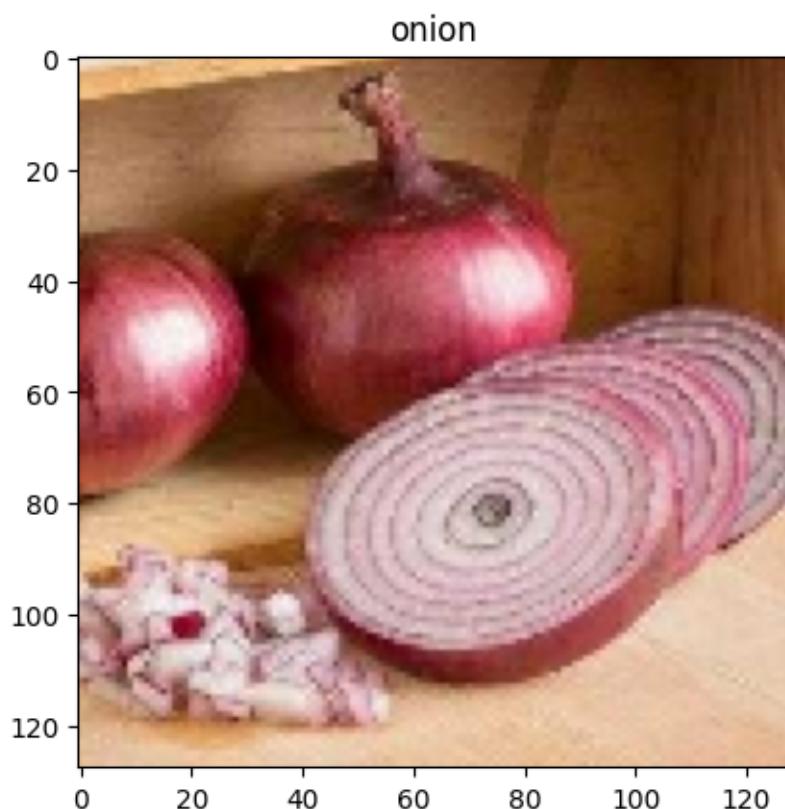
tomato



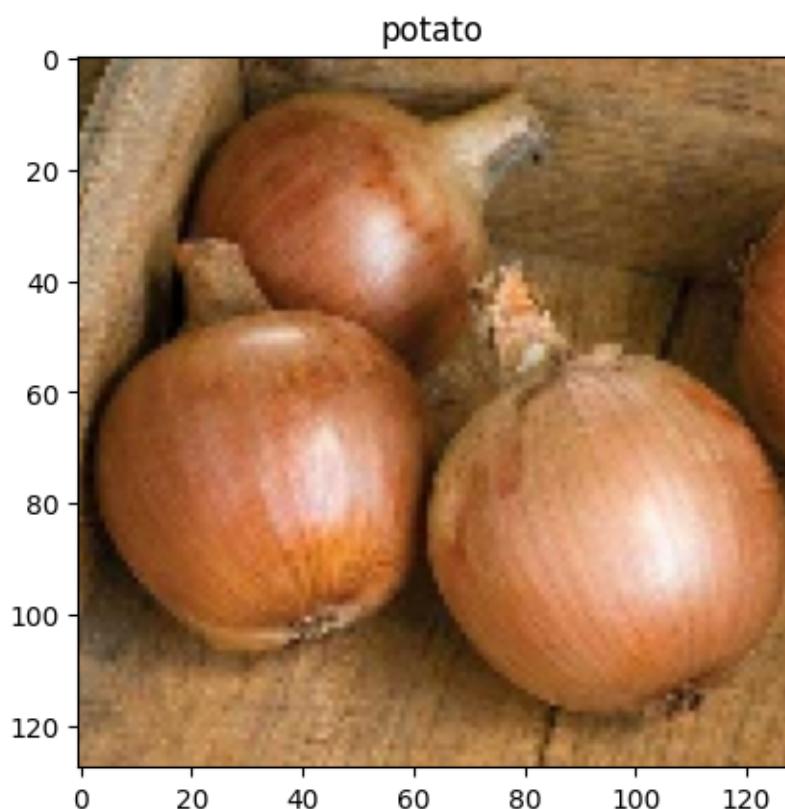
1/1 [=====] - 0s 34ms/step



1/1 [=====] - 0s 31ms/step



1/1 [=====] - 0s 40ms/step



1/1 [=====] - 0s 40ms/step



9 Applying Transfer Learning Method

```
[49]: base_model = tf.keras.applications.  
      ↪ResNet50(input_shape=(128,128,3),include_top=False,weights='imagenet')  
  
for layers in base_model.layers[-1:-1000:-1]:  
    #print(layers, layers.trainable)  
    layers.trainable = False  
  
inputs = tf.keras.Input(shape=(128,128,3))  
x = base_model(inputs,training=False)  
x = tf.keras.layers.GlobalAveragePooling2D()(x)  
outputs = tf.keras.layers.Dense(4,activation='softmax')(x)  
  
n_model = tf.keras.Model(inputs,outputs)
```

```
[50]: # for layers in base_model.layers:  
#       if(layers.trainable == False):  
#           print(layers, layers.trainable)
```

```
[51]: # for layers in base_model.layers:  
#     print(layers, layers.trainable)
```

```
[52]: n_model.summary()
```

```
Model: "model"  
-----  
Layer (type)          Output Shape         Param #  
=====-----  
input_2 (InputLayer)    [(None, 128, 128, 3)]   0  
resnet50 (Functional)  (None, 4, 4, 2048)      23587712  
global_average_pooling2d (G (None, 2048)  
lobalAveragePooling2D)  0  
dense (Dense)          (None, 4)             8196  
-----  
Total params: 23,595,908  
Trainable params: 8,196  
Non-trainable params: 23,587,712
```

```
[53]: n_model.compile(optimizer=tf.keras.optimizers.Adam(),  
                      loss='sparse_categorical_crossentropy',  
                      metrics=['accuracy'])  
  
filepath = "C://Users//gaura//Documents//GitHub//Scaler-Projects//Computer  
↪Vision//Project//TL_MODEL"  
callbacks = [  
    tf.keras.callbacks.ReduceLROnPlateau(  
        monitor="val_loss", factor=0.3, patience=5, min_lr=0.00001  
    ),  
    tf.keras.callbacks.ModelCheckpoint(filepath, save_weights_only=True,  
↪monitor='val_accuracy', mode='max', save_best_only=True),  
    tf.keras.callbacks.EarlyStopping(  
        monitor="val_loss", patience=10, min_delta=0.001, mode='min'  
    )  
]
```

```
[54]: tf.keras.backend.clear_session()
```

```
[55]: t_hist = n_model.  
↪fit(train_ds, validation_data=val_data, verbose=1, epochs=100, batch_size=32, callbacks=callback
```

Epoch 1/100

```
34/34 [=====] - 14s 247ms/step - loss: 1.3703 -  
accuracy: 0.2855 - val_loss: 1.3536 - val_accuracy: 0.2865 - lr: 0.0010  
Epoch 2/100  
34/34 [=====] - 8s 210ms/step - loss: 1.3423 -  
accuracy: 0.3131 - val_loss: 1.3322 - val_accuracy: 0.3729 - lr: 0.0010  
Epoch 3/100  
34/34 [=====] - 7s 186ms/step - loss: 1.3207 -  
accuracy: 0.3756 - val_loss: 1.3173 - val_accuracy: 0.3646 - lr: 0.0010  
Epoch 4/100  
34/34 [=====] - 8s 196ms/step - loss: 1.3010 -  
accuracy: 0.3825 - val_loss: 1.3040 - val_accuracy: 0.3625 - lr: 0.0010  
Epoch 5/100  
34/34 [=====] - 9s 218ms/step - loss: 1.2878 -  
accuracy: 0.4014 - val_loss: 1.2867 - val_accuracy: 0.3760 - lr: 0.0010  
Epoch 6/100  
34/34 [=====] - 9s 224ms/step - loss: 1.2731 -  
accuracy: 0.4230 - val_loss: 1.2723 - val_accuracy: 0.3833 - lr: 0.0010  
Epoch 7/100  
34/34 [=====] - 8s 214ms/step - loss: 1.2589 -  
accuracy: 0.4299 - val_loss: 1.2561 - val_accuracy: 0.4573 - lr: 0.0010  
Epoch 8/100  
34/34 [=====] - 7s 185ms/step - loss: 1.2417 -  
accuracy: 0.4667 - val_loss: 1.2540 - val_accuracy: 0.4354 - lr: 0.0010  
Epoch 9/100  
34/34 [=====] - 8s 196ms/step - loss: 1.2317 -  
accuracy: 0.4749 - val_loss: 1.2362 - val_accuracy: 0.4490 - lr: 0.0010  
Epoch 10/100  
34/34 [=====] - 8s 197ms/step - loss: 1.2273 -  
accuracy: 0.4464 - val_loss: 1.2383 - val_accuracy: 0.4292 - lr: 0.0010  
Epoch 11/100  
34/34 [=====] - 8s 195ms/step - loss: 1.2122 -  
accuracy: 0.4892 - val_loss: 1.2097 - val_accuracy: 0.4521 - lr: 0.0010  
Epoch 12/100  
34/34 [=====] - 9s 227ms/step - loss: 1.2027 -  
accuracy: 0.5076 - val_loss: 1.1983 - val_accuracy: 0.4833 - lr: 0.0010  
Epoch 13/100  
34/34 [=====] - 8s 190ms/step - loss: 1.1924 -  
accuracy: 0.5053 - val_loss: 1.2043 - val_accuracy: 0.4708 - lr: 0.0010  
Epoch 14/100  
34/34 [=====] - 12s 321ms/step - loss: 1.1797 -  
accuracy: 0.5159 - val_loss: 1.1924 - val_accuracy: 0.4875 - lr: 0.0010  
Epoch 15/100  
34/34 [=====] - 8s 180ms/step - loss: 1.1777 -  
accuracy: 0.5076 - val_loss: 1.1900 - val_accuracy: 0.4812 - lr: 0.0010  
Epoch 16/100  
34/34 [=====] - 9s 223ms/step - loss: 1.1714 -  
accuracy: 0.4970 - val_loss: 1.1774 - val_accuracy: 0.5052 - lr: 0.0010  
Epoch 17/100
```

```
34/34 [=====] - 8s 216ms/step - loss: 1.1619 -  
accuracy: 0.5168 - val_loss: 1.1684 - val_accuracy: 0.5104 - lr: 0.0010  
Epoch 18/100  
34/34 [=====] - 8s 195ms/step - loss: 1.1542 -  
accuracy: 0.5333 - val_loss: 1.1629 - val_accuracy: 0.5000 - lr: 0.0010  
Epoch 19/100  
34/34 [=====] - 8s 196ms/step - loss: 1.1478 -  
accuracy: 0.5255 - val_loss: 1.1693 - val_accuracy: 0.4792 - lr: 0.0010  
Epoch 20/100  
34/34 [=====] - 8s 193ms/step - loss: 1.1366 -  
accuracy: 0.5343 - val_loss: 1.1596 - val_accuracy: 0.4927 - lr: 0.0010  
Epoch 21/100  
34/34 [=====] - 9s 223ms/step - loss: 1.1295 -  
accuracy: 0.5434 - val_loss: 1.1515 - val_accuracy: 0.5260 - lr: 0.0010  
Epoch 22/100  
34/34 [=====] - 8s 214ms/step - loss: 1.1251 -  
accuracy: 0.5554 - val_loss: 1.1380 - val_accuracy: 0.5281 - lr: 0.0010  
Epoch 23/100  
34/34 [=====] - 8s 189ms/step - loss: 1.1263 -  
accuracy: 0.5490 - val_loss: 1.1341 - val_accuracy: 0.5156 - lr: 0.0010  
Epoch 24/100  
34/34 [=====] - 8s 195ms/step - loss: 1.1171 -  
accuracy: 0.5586 - val_loss: 1.1309 - val_accuracy: 0.5167 - lr: 0.0010  
Epoch 25/100  
34/34 [=====] - 8s 193ms/step - loss: 1.1096 -  
accuracy: 0.5582 - val_loss: 1.1187 - val_accuracy: 0.5188 - lr: 0.0010  
Epoch 26/100  
34/34 [=====] - 9s 226ms/step - loss: 1.1067 -  
accuracy: 0.5453 - val_loss: 1.1226 - val_accuracy: 0.5375 - lr: 0.0010  
Epoch 27/100  
34/34 [=====] - 8s 189ms/step - loss: 1.0959 -  
accuracy: 0.5683 - val_loss: 1.1190 - val_accuracy: 0.5292 - lr: 0.0010  
Epoch 28/100  
34/34 [=====] - 8s 196ms/step - loss: 1.1077 -  
accuracy: 0.5379 - val_loss: 1.1167 - val_accuracy: 0.5271 - lr: 0.0010  
Epoch 29/100  
34/34 [=====] - 8s 197ms/step - loss: 1.0888 -  
accuracy: 0.5678 - val_loss: 1.1130 - val_accuracy: 0.5281 - lr: 0.0010  
Epoch 30/100  
34/34 [=====] - 8s 208ms/step - loss: 1.0835 -  
accuracy: 0.5623 - val_loss: 1.1104 - val_accuracy: 0.5198 - lr: 0.0010  
Epoch 31/100  
34/34 [=====] - 8s 199ms/step - loss: 1.0836 -  
accuracy: 0.5724 - val_loss: 1.1021 - val_accuracy: 0.5344 - lr: 0.0010  
Epoch 32/100  
34/34 [=====] - 8s 204ms/step - loss: 1.0717 -  
accuracy: 0.5807 - val_loss: 1.1107 - val_accuracy: 0.5177 - lr: 0.0010  
Epoch 33/100
```

```
34/34 [=====] - 9s 224ms/step - loss: 1.0689 -  
accuracy: 0.5747 - val_loss: 1.0952 - val_accuracy: 0.5500 - lr: 0.0010  
Epoch 34/100  
34/34 [=====] - 8s 204ms/step - loss: 1.0672 -  
accuracy: 0.5802 - val_loss: 1.0988 - val_accuracy: 0.5312 - lr: 0.0010  
Epoch 35/100  
34/34 [=====] - 9s 232ms/step - loss: 1.0651 -  
accuracy: 0.5775 - val_loss: 1.0749 - val_accuracy: 0.5531 - lr: 0.0010  
Epoch 36/100  
34/34 [=====] - 9s 215ms/step - loss: 1.0631 -  
accuracy: 0.5816 - val_loss: 1.0812 - val_accuracy: 0.5396 - lr: 0.0010  
Epoch 37/100  
34/34 [=====] - 9s 223ms/step - loss: 1.0589 -  
accuracy: 0.5743 - val_loss: 1.0912 - val_accuracy: 0.5312 - lr: 0.0010  
Epoch 38/100  
34/34 [=====] - 6s 148ms/step - loss: 1.0494 -  
accuracy: 0.5793 - val_loss: 1.0860 - val_accuracy: 0.5375 - lr: 0.0010  
Epoch 39/100  
34/34 [=====] - 6s 144ms/step - loss: 1.0506 -  
accuracy: 0.5802 - val_loss: 1.0759 - val_accuracy: 0.5500 - lr: 0.0010  
Epoch 40/100  
34/34 [=====] - 7s 166ms/step - loss: 1.0424 -  
accuracy: 0.5867 - val_loss: 1.0761 - val_accuracy: 0.5635 - lr: 0.0010  
Epoch 41/100  
34/34 [=====] - 6s 136ms/step - loss: 1.0387 -  
accuracy: 0.5940 - val_loss: 1.0767 - val_accuracy: 0.5469 - lr: 3.0000e-04  
Epoch 42/100  
34/34 [=====] - 6s 132ms/step - loss: 1.0383 -  
accuracy: 0.5954 - val_loss: 1.0771 - val_accuracy: 0.5479 - lr: 3.0000e-04  
Epoch 43/100  
34/34 [=====] - 6s 135ms/step - loss: 1.0358 -  
accuracy: 0.5991 - val_loss: 1.0655 - val_accuracy: 0.5573 - lr: 3.0000e-04  
Epoch 44/100  
34/34 [=====] - 5s 133ms/step - loss: 1.0418 -  
accuracy: 0.5913 - val_loss: 1.0743 - val_accuracy: 0.5490 - lr: 3.0000e-04  
Epoch 45/100  
34/34 [=====] - 6s 148ms/step - loss: 1.0396 -  
accuracy: 0.5903 - val_loss: 1.0703 - val_accuracy: 0.5708 - lr: 3.0000e-04  
Epoch 46/100  
34/34 [=====] - 6s 138ms/step - loss: 1.0353 -  
accuracy: 0.5977 - val_loss: 1.0683 - val_accuracy: 0.5510 - lr: 3.0000e-04  
Epoch 47/100  
34/34 [=====] - 6s 152ms/step - loss: 1.0388 -  
accuracy: 0.5908 - val_loss: 1.0564 - val_accuracy: 0.5792 - lr: 3.0000e-04  
Epoch 48/100  
34/34 [=====] - 6s 136ms/step - loss: 1.0348 -  
accuracy: 0.5949 - val_loss: 1.0747 - val_accuracy: 0.5437 - lr: 3.0000e-04  
Epoch 49/100
```

```
34/34 [=====] - 6s 134ms/step - loss: 1.0331 -
accuracy: 0.5926 - val_loss: 1.0637 - val_accuracy: 0.5583 - lr: 3.0000e-04
Epoch 50/100
34/34 [=====] - 6s 135ms/step - loss: 1.0266 -
accuracy: 0.6014 - val_loss: 1.0674 - val_accuracy: 0.5698 - lr: 3.0000e-04
Epoch 51/100
34/34 [=====] - 6s 133ms/step - loss: 1.0387 -
accuracy: 0.5899 - val_loss: 1.0623 - val_accuracy: 0.5625 - lr: 3.0000e-04
Epoch 52/100
34/34 [=====] - 6s 137ms/step - loss: 1.0367 -
accuracy: 0.5931 - val_loss: 1.0768 - val_accuracy: 0.5510 - lr: 3.0000e-04
Epoch 53/100
34/34 [=====] - 6s 141ms/step - loss: 1.0243 -
accuracy: 0.5982 - val_loss: 1.0484 - val_accuracy: 0.5781 - lr: 9.0000e-05
Epoch 54/100
34/34 [=====] - 6s 142ms/step - loss: 1.0274 -
accuracy: 0.5926 - val_loss: 1.0583 - val_accuracy: 0.5635 - lr: 9.0000e-05
Epoch 55/100
34/34 [=====] - 6s 134ms/step - loss: 1.0221 -
accuracy: 0.5986 - val_loss: 1.0615 - val_accuracy: 0.5615 - lr: 9.0000e-05
Epoch 56/100
34/34 [=====] - 6s 132ms/step - loss: 1.0331 -
accuracy: 0.5917 - val_loss: 1.0691 - val_accuracy: 0.5594 - lr: 9.0000e-05
Epoch 57/100
34/34 [=====] - 6s 134ms/step - loss: 1.0313 -
accuracy: 0.5926 - val_loss: 1.0703 - val_accuracy: 0.5531 - lr: 9.0000e-05
Epoch 58/100
34/34 [=====] - 6s 136ms/step - loss: 1.0267 -
accuracy: 0.5945 - val_loss: 1.0573 - val_accuracy: 0.5698 - lr: 9.0000e-05
Epoch 59/100
34/34 [=====] - 6s 138ms/step - loss: 1.0288 -
accuracy: 0.5926 - val_loss: 1.0700 - val_accuracy: 0.5500 - lr: 2.7000e-05
Epoch 60/100
34/34 [=====] - 6s 132ms/step - loss: 1.0277 -
accuracy: 0.5963 - val_loss: 1.0611 - val_accuracy: 0.5635 - lr: 2.7000e-05
Epoch 61/100
34/34 [=====] - 6s 136ms/step - loss: 1.0253 -
accuracy: 0.5926 - val_loss: 1.0610 - val_accuracy: 0.5615 - lr: 2.7000e-05
Epoch 62/100
34/34 [=====] - 6s 135ms/step - loss: 1.0336 -
accuracy: 0.5899 - val_loss: 1.0579 - val_accuracy: 0.5625 - lr: 2.7000e-05
Epoch 63/100
34/34 [=====] - 6s 134ms/step - loss: 1.0277 -
accuracy: 0.5995 - val_loss: 1.0613 - val_accuracy: 0.5615 - lr: 2.7000e-05
```

```
[64]: # n_model.save('tl_model.pb')
```

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,

```
_jit_compiled_convolution_op, _jit_compiled_convolution_op,  
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing  
5 of 53). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: tl_model.pb\assets
```

```
INFO:tensorflow:Assets written to: tl_model.pb\assets
```

```
[65]: load_model = tf.keras.models.load_model("C://Users//gaura//Documents//GitHub//  
Scaler-Projects//Computer Vision//Project//Notebook//tl_model.pb")
```

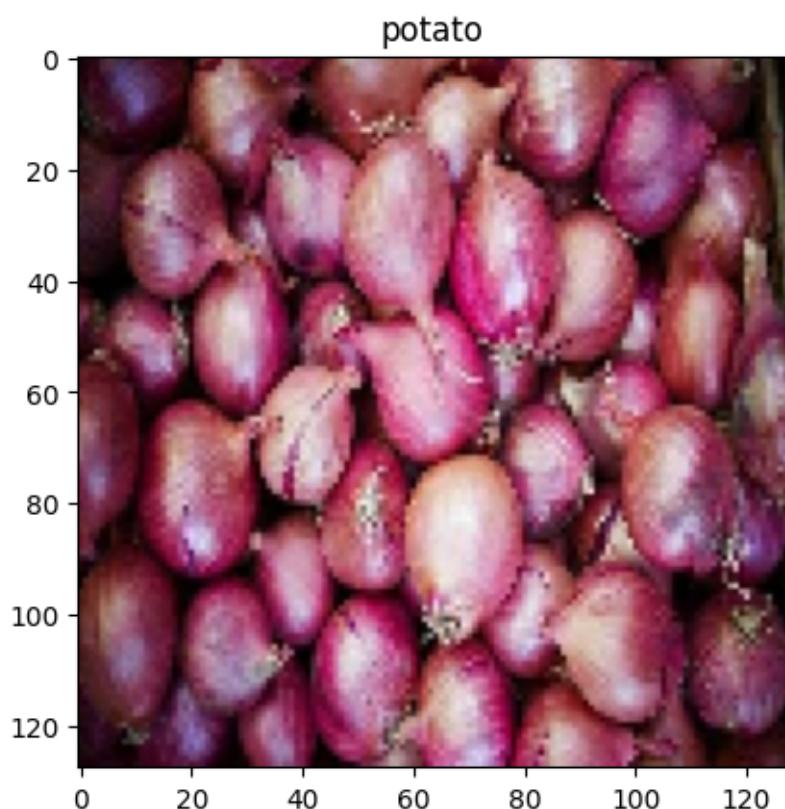
```
[66]: # Evaluate the model  
loss, acc = load_model.evaluate(test_data, verbose=2)  
print("Restored model, accuracy: {:.2f}%".format(100 * acc))
```

```
6/6 - 2s - loss: 1.1665 - accuracy: 0.4644 - 2s/epoch - 338ms/step  
Restored model, accuracy: 46.44%
```

```
[67]: images,label = next(iter(test_data))
```

```
[68]: for img in images.numpy()[:5]:  
    #print(img.shape)  
    temp = (np.expand_dims(img, axis=0))  
    pred = load_model.predict(temp)  
    labels = reference_mapping_label[np.argmax(pred)]  
    plt.imshow(img)  
    plt.title(labels)  
    plt.show()
```

```
1/1 [=====] - 1s 1s/step
```

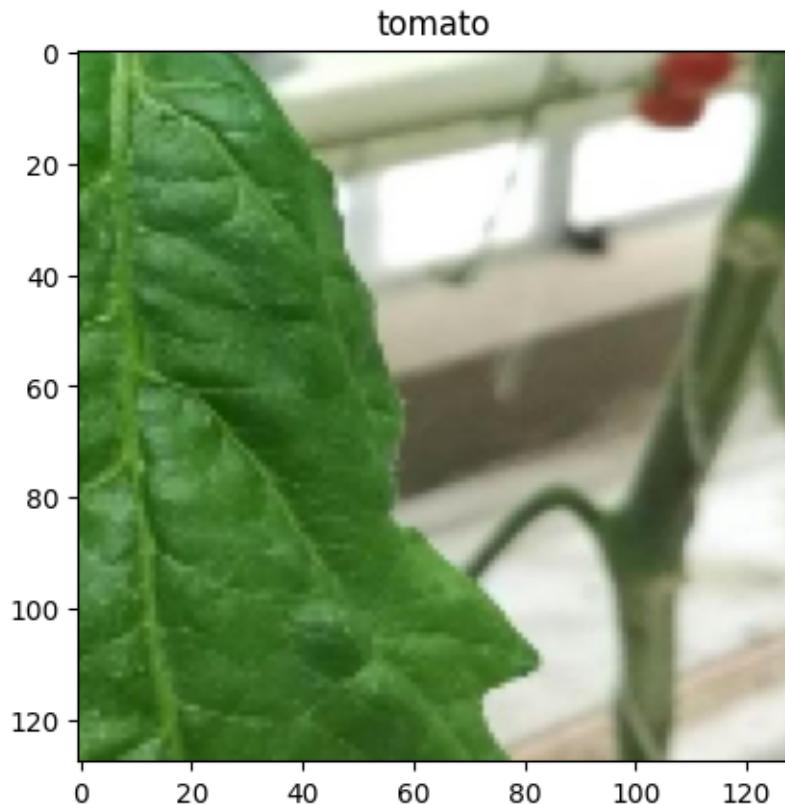


1/1 [=====] - 0s 33ms/step

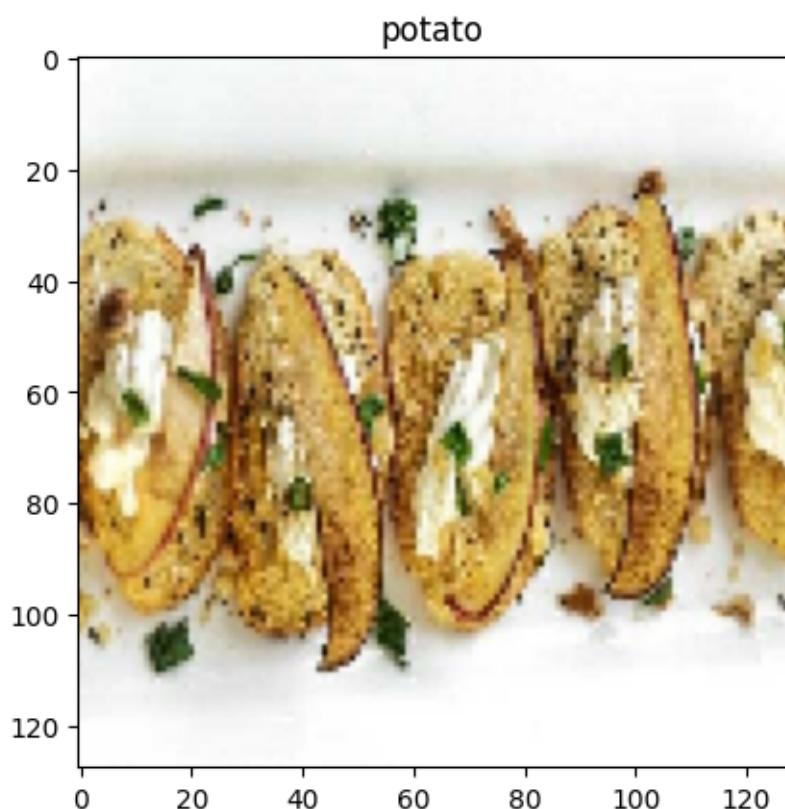
indian market



1/1 [=====] - 0s 24ms/step



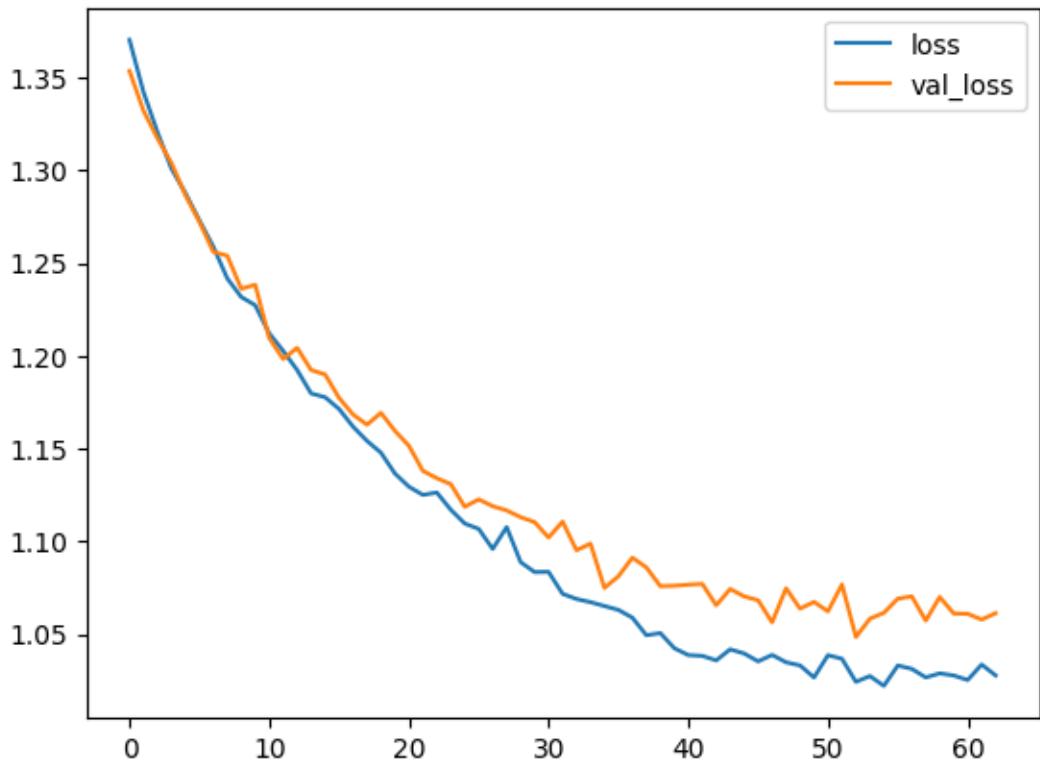
1/1 [=====] - 0s 26ms/step



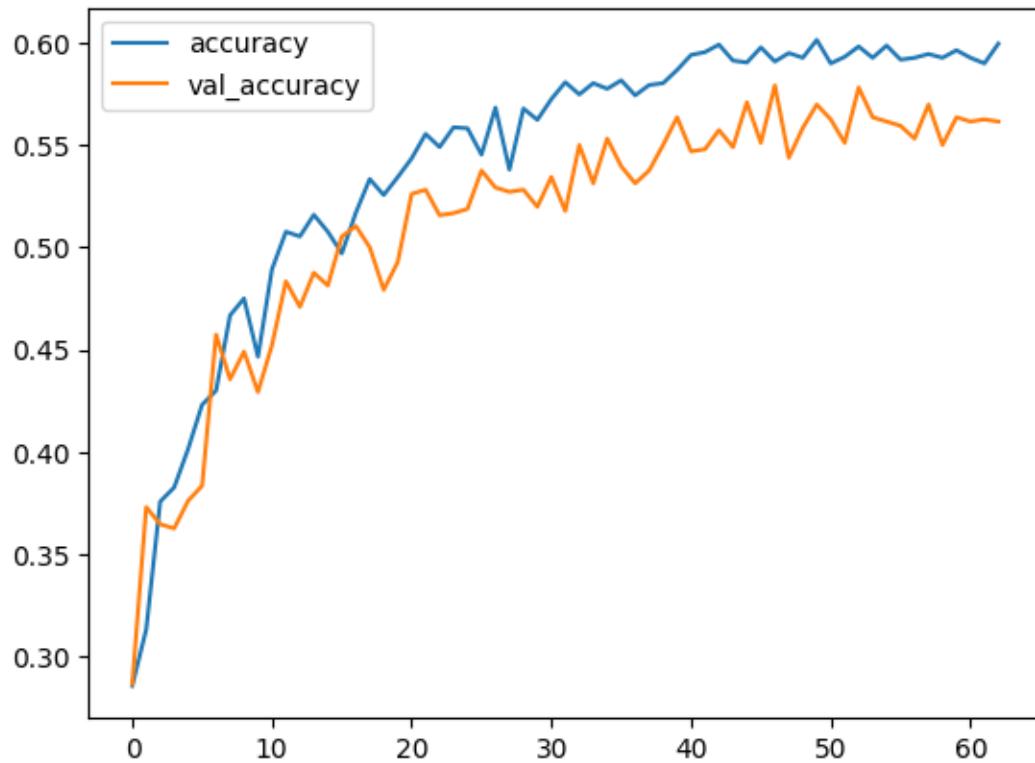
1/1 [=====] - 0s 32ms/step



```
[69]: plt.plot(t_hist.history['loss'],label='loss')
plt.plot(t_hist.history['val_loss'],label='val_loss')
plt.legend()
plt.show()
```



```
[70]: plt.plot(t_hist.history['accuracy'],label='accuracy')
plt.plot(t_hist.history['val_accuracy'],label='val_accuracy')
plt.legend()
plt.show()
```



10 Summary of Model

1. Both the model have their own way of generalization.
2. CNN from scratch has achieved accuracy of 75.21% but was unable to learn the params and end up with missclassification.
3. Using Transfer learning the model has achieved accuracy of 46% accuracy on test but has less missclassification problems.

[]: